

ATP

Account Technology Profiler

User Guide

Your AI assistant for running your accounts — powered by a local NVIDIA model
Nemotron-Super-49B on 2× H100 | private & local | no API keys

SHI International Corp. | Solutions Engineering
Version 2.0 | June 2026

CONTENTS

- 1 Welcome To ATP
- 2 Starting Up & The Pre-Demo Check
- 3 The Home Screen & Command Box
- 4 Asking About Your Customers (RAG)
- 5 Researching A Technology (Web)
- 6 “Is It A Good Fit?” — Brainwaves
- 7 Briefings & Other Skills
- 8 SEdraw — Network Diagrams
- 9 SEdraw — Convert From A Word Doc
- 10 SEdraw — Vision Boards
- 11 The Customer Vault & Profiles
- 12 Artifacts & Brainwave History
- 13 Deleting Things Safely
- 14 Guardrails — Keeping ATP On-Task
- 15 Settings
- 16 Tips & Recovery

1. WELCOME TO ATP

ATP — the Account Technology Profiler — is your AI assistant for running your accounts. You type what you want in one command box and ATP figures out what to do: answer questions about your customers, research a technology on the web, judge whether a product fits an account, build briefings, or draw network and vision diagrams.

Everything runs on a **local NVIDIA Nemotron model on two H100 GPUs** in your own environment — **no cloud AI service, no API keys, and your customer data never leaves the machine** (the one exception is the web-research feature, which searches the public web when you ask it to).

This guide walks through everything ATP can do and how to use it.

2. STARTING UP & THE PRE-DEMO CHECK

Open the **Account Technology Profiler** shortcut on the desktop to launch the app window. To confirm the AI model is live before a demo, run the health check in a terminal:

```
./demo_health.sh
```

It prints a green board: both H100 GPUs, the Nemotron model loaded across them, the inference engine, a live test generation (with speed), and the app. If something is down, run:

```
./demo_health.sh --recover      # restarts the model only if it's actually down
```

- All green → you're ready to demo.
- Red items → the recovery line above fixes the model; the board also lists the exact commands.

3. THE HOME SCREEN & COMMAND BOX

The Home screen centers on the **command box** — type anything and press Enter. Underneath are two starter chips you can click to pre-fill a request:

- ■ **Create a strategic account briefing for a customer**
- ■ **Research whether a new technology fits a customer's infrastructure**

Click a chip and the first [bracketed] part is selected — just type the customer or technology over it and press Enter. Around the box you'll also find your live tasks (**Working on**, with an X to cancel), **Recent Artifacts**, and buttons for the Customer Vault, Skills, Cron Jobs, Brainwave History, and Artifacts.

Whatever you type is understood and routed automatically — you don't pick a mode. Examples:

Type this...	ATP does...
which customers use Palo Alto firewalls?	Searches every customer and answers with citations
does Aspen Health use CrowdStrike?	Checks that one customer; shows where it found it
research the latest Palo Alto firewalls	Searches the web and writes a sourced brief
is Verkada a good fit for Aspen Health?	Runs a multi-agent analysis + a verdict
strategic briefing for Aspen Health	Builds a full strategic briefing

4. ASKING ABOUT YOUR CUSTOMERS (RAG)

Ask a plain question about your accounts and ATP searches the Customer Vault — trip reports, the technology profile, vision/network diagrams, and filed artifacts — then answers with **citations** to exactly where each fact came from.

- **Across all customers:** “which of my customers use CrowdStrike?” → a list, each with where it was found.
- **One customer:** “does Aspen Health use Palo Alto?” → yes/no with the specific trip reports / profile entries.
- If nothing matches, ATP says so plainly rather than guessing.

5. RESEARCHING A TECHNOLOGY (WEB)

Ask ATP to **research** something and it goes to the web (not your vault): “research the latest Palo Alto firewalls”, “what's new with Zscaler”. A ■ research task runs, then a sourced brief appears as an artifact and a Brainwave-History entry so you can see the reasoning.

This is the only feature that reaches the internet, and only the search words leave the machine. You can block specific sources (see Guardrails).

6. “IS IT A GOOD FIT?” — BRAINWAVES

Ask whether a product fits a customer — “is Verkada a good fit for Aspen Health?” — and ATP runs a **multi-agent Brainwave**:

- ■ one agent researches the product on the web, while ■ another studies that customer's saved infrastructure — **at the same time**;
- ■ a third agent merges both into a grounded verdict (fit, where it fits, risks, recommended next step);
- ■ a fourth agent **fact-checks the verdict** against the evidence and flags anything unsupported.

The result lands in **Brainwave History** as a full receipt: what it decided and why, which agents ran (and that they ran in parallel on the Nemotron model), the verdict, and the self-check. This is where you go to see the logic behind any answer.

7. BRIEFINGS & OTHER SKILLS

Type “strategic briefing for <customer>” (or use the chip) and ATP builds a full briefing — executive summary, initiatives, opportunities, risks, talking points, and a decision-tree diagram — grounded in that customer's saved data. Other Brain skills (account snapshot, battlecard, whitespace, follow-up, environment topology) run the same way from the command box. The deliverable is saved to Artifacts and can be filed to the customer.

8. SEDRAW — NETWORK DIAGRAMS

Open **Skills** → **SEdraw Network Diagrams** and pick a customer. You can:

- **Generate** a network map from an Excel intake already in that customer's folder;
- **Create New** — start from a blank template and fill it in the built-in editor;
- **Edit** an existing intake in the GUI editor, then regenerate;
- **Delete** a diagram (PIN-guarded — see Deleting Things).

The generated diagram opens right in the window; a download button gives you the .drawio file.

9. SEDRAW — CONVERT FROM A WORD DOC

In the same Network view, click **Convert from DOCX** and pick a Word intake (.docx). The Nemotron model reads the document — paragraphs and tables — and does its best to map it onto the network intake form, then generates the map automatically.

- It maps what it reasonably can (sites, firewalls, switches, servers, storage, connectivity) and skips anything that doesn't fit — it won't invent devices.
- The mapped intake is saved as an editable spreadsheet, so you can open the editor, fix anything, and regenerate.
- It runs in the background (you'll see progress); large documents are summarized to fit the model.

10. SEDRAW — VISION BOARDS

Open **Skills** → **SEdraw Vision Boards** and pick a customer (the Vision Studio opens). Describe the environment in plain English in two boxes — **Current State** and **Future State** — and click Generate. Nemotron builds one combined architecture diagram (current on top, future below).

- It runs as a background task; the diagram appears and is filed as an artifact.
- **History** lists earlier boards — open or edit any of them in the full editor.
- **Delete** a board (PIN-guarded).

11. THE CUSTOMER VAULT & PROFILES

Customer Vault lists every account. Open one to see its trip reports, its **Account Technology Profile** (technologies, overview, key contacts, topology), network/vision diagrams, and artifacts.

- **Generate** rebuilds the technology profile from the customer's trip reports; each run is saved as a dated **generation** snapshot you can compare over time.
- Your hand-edits are preserved across regenerations (the system merges rather than overwrites).
- **Topology** shows the rendered network/technology diagram; older generations each have a **X** to remove (the live 'Current' profile is protected).

12. ARTIFACTS & BRAINWAVE HISTORY

- **Artifacts** — every deliverable ATP produces (briefings, diagrams, research briefs, profiles). Search by title/customer/content, open them in-app, and file them to a customer.
- **Brainwave History** — a reasoning receipt for every command: what ATP decided and why, which agents ran (on the Nemotron model), the rationale, and (for brainwaves) the verdict + self-check. This is your audit trail of the AI's thinking.

13. DELETING THINGS SAFELY

Every delete in ATP — a customer, an artifact, a vision board, a network diagram, or a profile generation — is a **soft delete**. Clicking  asks for your **6-digit System PIN**; the item then moves to a recycle bin and disappears from the app, but is **never erased** — it can be restored from `vault/recyclebin/`.

- Set your PIN in **Settings** first; without it, deletion is blocked (a safety guard).
- This is the same safe-delete behavior everywhere in the app.

14. GUARDRAILS — KEEPING ATP ON-TASK

ATP can run a **guardrail** on the command box that keeps it on-topic (your customers + technology research) and politely declines off-topic, unsafe, or prompt-injection requests. It is **off by default** and **fail-open** (it never blocks legitimate work if something goes wrong).

Turning it on

- Quick terminal toggle: `./guardrails.sh on` (or `off` / `status`).
- Or use the **Guardrails Control Panel** desktop shortcut (opens at `http://localhost:8090`) — a separate window that can't affect the main app.

What you can control (the levers)

- Master on/off · keep on-topic · block unsafe · block prompt-injection · strictness (lenient / balanced / strict).
- **Blocked words** — refuse any command containing them (e.g. "reddit").
- **Blocked research sources** — stop the web-research agent from using certain domains (e.g. `reddit.com`).
- **Live Test box** — type a phrase to see ALLOW or REFUSE with the reason, before you enable it.
- **Recent blocks** — a live log of what got declined. Click  **Save** to confirm changes (they take effect on your next command — no restart).

15. SETTINGS

- **System PIN** — required to delete anything.
- **Timezone** — used by scheduled jobs.
- **Guardrails** — enable/disable the command-box rail.
- **Brain Rules** — guidance injected into the assistant's prompts (rule #1: look at the customer's saved data first).

16. TIPS & RECOVERY

If...	Do this
You want to prove the model is live	<code>./demo_health.sh</code> — shows the 2× H100s, the model, and a live generation.
The model seems down mid-demo	<code>./demo_health.sh --recover</code> (heals the model only if it's actually down).
A task is taking too long	Click the X on its chip in 'Working on' — it stops the task and any sub-agents.
The guardrail blocked something you wanted	Lower strictness / clear blocked words in the panel, or <code>./guardrails.sh off</code> .
A Word→network convert looks thin	Open the generated intake in the editor, fix it, and regenerate.
You deleted something by mistake	It's in <code>vault/recyclebin/</code> — nothing is ever hard-deleted.